PAPER GRANNY
arXiv Paper Interpretation Agent

ReAct: Synergizing Reasoning and Acting in Language Models

Paper Interpretation Report

Date: October 2022

Paper Information

Original Paper: arXiv:2210.03629

PAPER GRANNY
arXiv Paper Interpretation Agent

目录

1 Background Knowledge **2**

2 Core Problem and Innovation **3**

3 Method Details **3**

3.1 The ReAct Framework 3

3.2 How ReAct Works in Practice 4

3.3 Types of Thoughts in ReAct 5

4 Experimental Results **6**

4.1 Knowledge-Intensive Reasoning Tasks 6

4.2 Decision Making Tasks 7

4.3 Human Interpretability 8

5 Summary **9**

6 References **9**

1 Background Knowledge

What are Large Language Models?

Imagine you have a very smart friend who has read almost everything on the internet. This friend can:

- **Answer questions** about many topics
- **Write stories** or explanations
- **Have conversations** that make sense

This is what a **Large Language Model (LLM)** does! It's been trained on huge amounts of text and can generate human-like responses.

Chain-of-Thought Reasoning

Sometimes, when you face a difficult problem, you think step by step before giving an answer. For example:

> "I need to find the capital of France. First, I know Paris is a city. Let me think... Yes, Paris is the capital of France."

Chain-of-Thought (CoT) prompting asks AI models to show their thinking process before giving an answer, which often leads to better results.

The Problem: Acting vs. Reasoning

Current AI models have two separate capabilities:

1. **Reasoning**: Thinking about problems step by step (like solving a math problem)
2. **Acting**: Taking actions in the real world (like searching the internet, clicking buttons)

The problem is that these two capabilities are studied separately. Humans naturally combine them—we think, then act, then think again based on what we learned!

2 Core Problem and Innovation

What is ReAct?

ReAct = **Re**asoning + **Act**ing
The key idea is simple: **Let AI models think AND act together, alternating between them**.

- **Reasoning**: Generate thoughts to plan, track progress, handle exceptions
- **Acting**: Take actions to interact with the environment (search Wikipedia, click buttons, etc.)
- **Together**: Reasoning guides actions, and actions provide information for better reasoning

Why is this Important?

Think about cooking a meal:

- You **think**: "I need to cut vegetables, then heat the pot"
- You **act**: Open the fridge, chop vegetables
- You **think again**: "Oh, I don't have salt, let me use soy sauce instead"
- You **act again**: Add soy sauce

This tight loop of thinking and acting is what makes humans good at solving problems. ReAct tries to give AI models the same ability!

3 Method Details

3.1 The ReAct Framework

The ReAct framework works by alternating between **thoughts** and **actions**:

1. **Receive observation**: Get information from the environment
2. **Generate thought**: Think about what to do next (this doesn't change the environment)
3. **Generate action**: Take a real action that affects the environment
4. **Get new observation**: See what happened after the action
5. **Repeat** until the task is complete

Key Formula: Augmented Action Space

$$\hat{\mathcal{A}} = \mathcal{A} \cup \mathcal{L}$$

Where:

- $\mathcal{A}$ = Original action space (e.g., search, click, buy)
- $\mathcal{L}$ = Language space (free-form thoughts)
- $\hat{\mathcal{A}}$ = Combined action space

Intuition: The model can either take a real action OR generate a thought. Thoughts help plan and reason, while actions interact with the world.

3.2 How ReAct Works in Practice

Let me show you a concrete example from the paper:

Example: Fact Verification

Task: Verify if "Bermuda Triangle is in the Pacific Ocean" is true.
ReAct trajectory:

1. **Thought 1**: "I need to search Bermuda Triangle and find if it is in the Pacific Ocean."
2. **Action 1**: Search[Bermuda Triangle]
3. **Observation 1**: "Reportedly haunted locations in the western part of the North Atlantic Ocean..."
4. **Thought 2**: "The observation says it's in the North Atlantic Ocean, not Pacific Ocean."
5. **Action 2**: Finish[REFUTES]

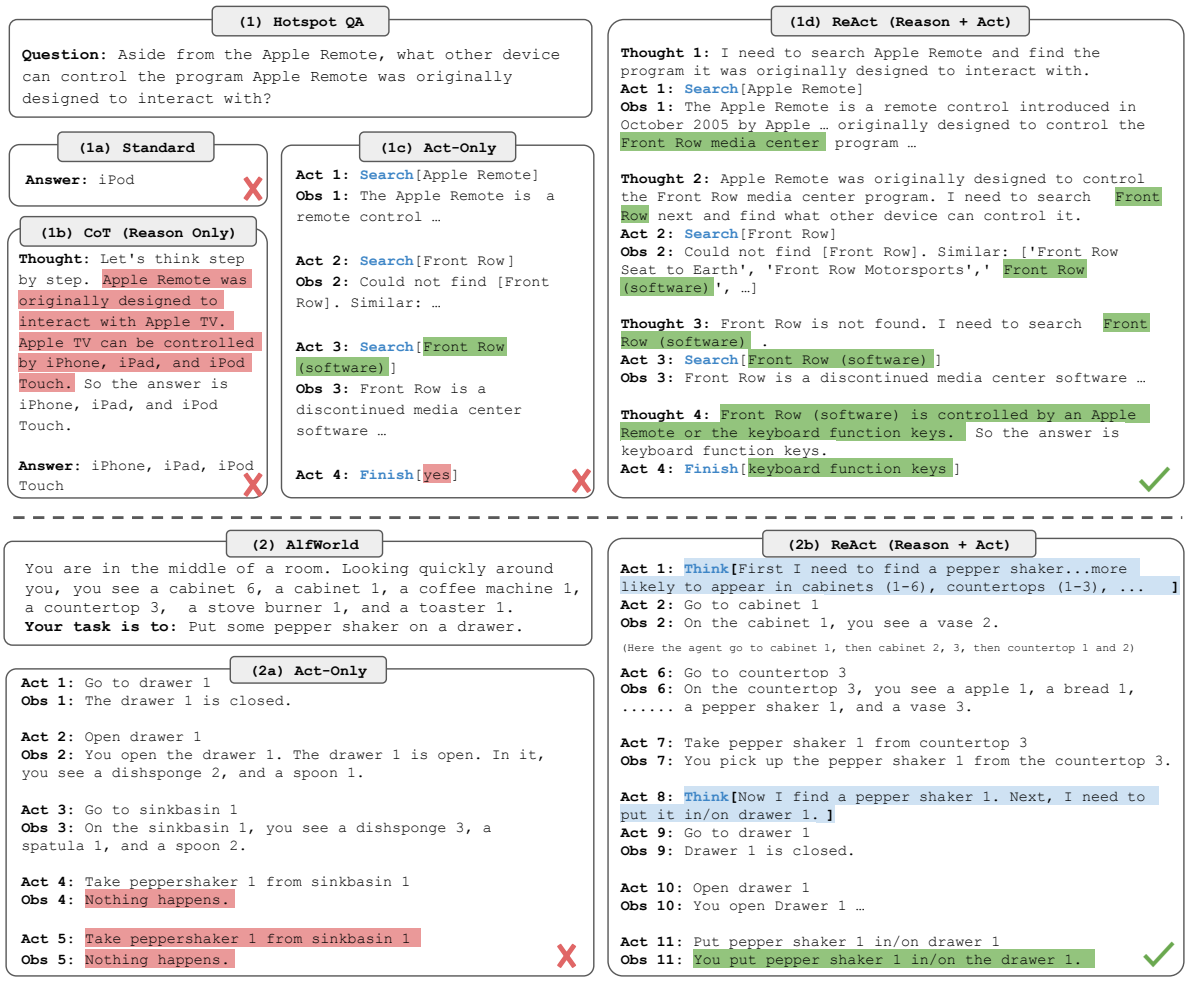

图 1: **Original Figure 1**: ReAct vs. other methods. (a) Standard prompting (no reasoning, no acting). (b) Chain-of-Thought (reasoning only, no acting). (c) Acting-only (acting without reasoning). (d) **ReAct** (alternating reasoning and acting).

Why ReAct is Better

- **More factual**: Can search Wikipedia to get accurate information
- **More interpretable**: Humans can read the thought process
- **More robust**: Can handle exceptions and adjust plans
- **Better for complex tasks**: Can decompose goals and track progress

3.3 Types of Thoughts in ReAct

The paper identifies several useful types of thoughts:

Types of Thoughts

1. **Goal decomposition**: Break down complex tasks into smaller steps
2. **Information extraction**: Pull important facts from observations
3. **Commonsense reasoning**: Use general knowledge to make inferences
4. **Progress tracking**: Keep track of what's been done and what's next
5. **Exception handling**: Adjust plans when things don't go as expected
6. **Answer synthesis**: Combine all information to form final answer

4 Experimental Results

4.1 Knowledge-Intensive Reasoning Tasks

The authors tested ReAct on two challenging tasks:

Task 1: HotPotQA - Multi-hop Question Answering

Goal: Answer questions that require information from multiple Wikipedia articles.
Challenge: Need to search, find relevant information, and combine it.

Task 2: FEVER - Fact Verification

Goal: Verify if claims are SUPPORTS, REFUTES, or NOT ENOUGH INFO based on Wikipedia.
Challenge: Need to search and compare claims with actual facts.

表 1: Performance on HotPotQA and FEVER (higher is better)

Method	HotPotQA	FEVER
Standard (no reasoning, no acting)	18.9	48.1
Chain-of-Thought (reasoning only)	29.4	56.3
Acting-only (acting without reasoning)	25.2	52.8
ReAct (ours)	**27.4**	**60.9**

Key Findings

- **ReAct beats Acting-only**: Shows reasoning helps guide actions
- **ReAct beats CoT on FEVER**: Accessing real information helps verification
- **CoT slightly better on HotPotQA**: Sometimes pure reasoning works better
- **Best approach**: Combine ReAct with CoT for different tasks

4.2 Decision Making Tasks

The authors also tested ReAct on interactive environments:

Task 1: ALFWorld - Text-based Game

Goal: Navigate a simulated house to complete tasks (e.g., "examine paper under desklamp").
Challenge: Need to plan, track progress, and use commonsense about where items might be.

Task 2: WebShop - Online Shopping

Goal: Purchase products based on user instructions (e.g., "nightstand with drawers, under $140").
Challenge: Navigate noisy product descriptions, make choices, and complete purchases.

表 2: Performance on ALFWorld and WebShop (higher is better)

Method	ALFWorld	WebShop
Imitation Learning (IL)	37.0	43.8
IL + Reinforcement Learning (RL)	48.0	53.1
Acting-only (1-shot)	45.0	52.5
ReAct (1-shot)	**71.0**	**63.5**

Amazing Results!

- **ALFWorld**: ReAct achieves 71% success rate vs. 48% for best baseline (34% absolute improvement!)
- **WebShop**: ReAct achieves 63.5% success rate vs. 53.1% for best baseline (10% improvement)
- **Few-shot learning**: ReAct uses only 1-2 examples, while baselines need thousands of training examples!

PAPER GRANNY
arXiv Paper Interpretation Agent

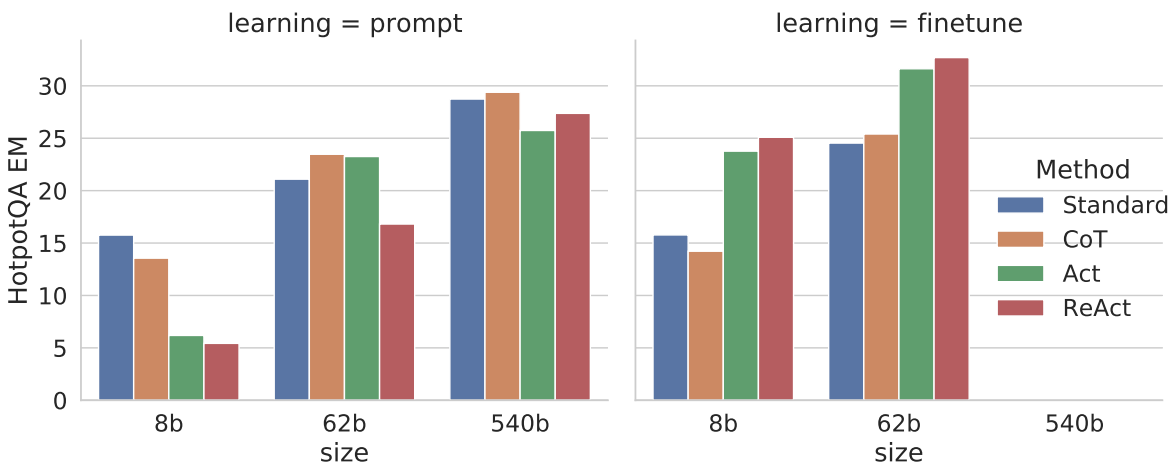

图 2: **Original Figure 3**: Fine-tuning results. With fine-tuning, ReAct becomes the best method across all model sizes.

4.3 Human Interpretability

One of ReAct's biggest advantages is that humans can understand and control the model's decisions:

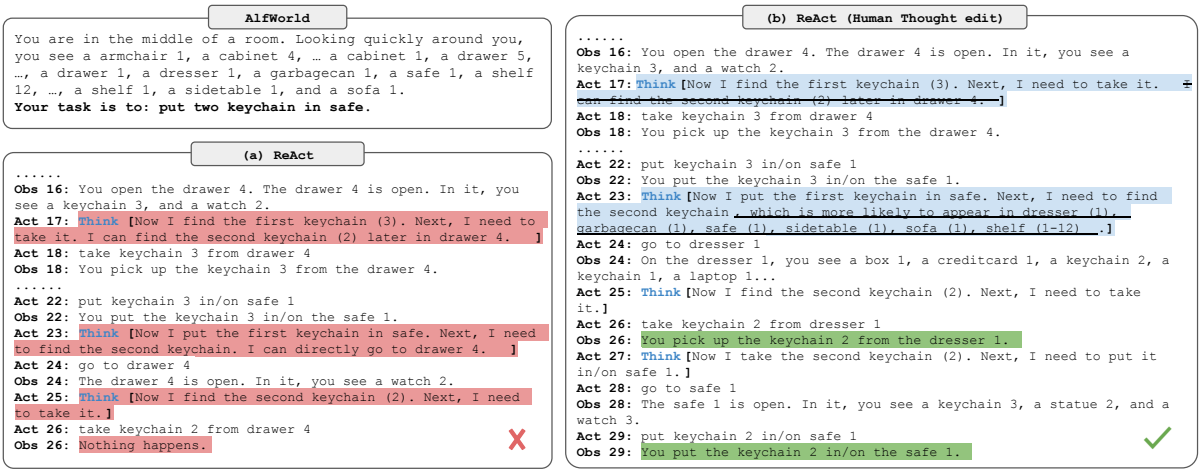

图 3: **Original Figure 4**: Human-in-the-loop editing. Humans can edit ReAct's thoughts to correct errors and guide the model.

Why This Matters

- **Transparency**: You can see exactly why the model made each decision
- **Control**: Humans can edit thoughts to fix errors
- **Trust**: Easier to trust decisions when you can understand them
- **Collaboration**: Humans and AI can work together more effectively

5 Summary

Key Contributions

1. **Introduce ReAct**: A new paradigm that combines reasoning and acting in language models
2. **Extensive experiments**: Show ReAct works well on diverse tasks (QA, fact verification, games, shopping)
3. **Performance boost**: Outperforms baselines by large margins with very few examples
4. **Interpretability**: Provides human-readable reasoning traces for better trust and control

Limitations

- **Limited reasoning/acting types**: Current implementation only supports specific reasoning and action patterns **Scaling challenge**: Complex tasks with large action spaces need more training examples
- **Fine-tuning needed for best performance**: Prompting alone may not be enough for very complex tasks
- **Potential risks**: Connecting LLMs to external environments could have safety concerns

Future Directions

- Scale up ReAct to more tasks and domains
- Combine with reinforcement learning for even better performance
- Add more types of reasoning and actions
- Improve fine-tuning with more human-written examples
- Explore human-AI collaboration more systematically

6 References

- Wei, J., et al. (2022). Chain-of-Thought Prompting Elicits Reasoning in Large Language Models.
- Shridhar, M., et al. (2020). ALFWorld: Aligning Text and Embodied Environments.
- Yao, S., et al. (2022). WebShop: Towards Realistic Commercial Web Navigation.
- Yang, Z., et al. (2018). HotpotQA: A Dataset for Diverse Question Answering with Multi-hop Reasoning.
- Thorne, J., et al. (2018). FEVER: A Large-scale Dataset for Fact Extraction and VERification.